

The background features a light blue gradient with decorative circuit-like lines in a darker blue. These lines, consisting of straight segments and small circles at junctions, are primarily concentrated on the left side and extend towards the top and bottom edges of the frame.

# A tutorial on efficient hash table algorithms

Yebangyu(bangyu.yby@alipay.com)

OceanBase

<http://www.yebangyu.org>

2015.12.29

# Outline

- Why Separate Chaining & Open Addressing are not enough
- Cuckoo Hashing
- Hopscotch Hashing

# Outline

- Why Separate Chaining & Open Addressing are not enough
- Cuckoo Hashing
- Hopscotch Hashing

# Disadvantages of Separate Chaining

- Poor cache performance
- $O(n)$  worst case behavior

# Disadvantages of Open Addressing

- Performance degrades as the table fills up
- Clustering of keys cause a large variance in performance

# Outline

- Why Separate Chaining & Linear Probe are not enough
- Cuckoo Hashing
- Hopscotch Hashing


# Cuckoo Hashing

- Motivation
- Intuition
- How it works
- Implementation details

# Authors of cuckoo hashing

- **Rasmus Pagh**, Professor of IT University of Copenhagen
- Research in algorithms for massive data, database performance
- describe cuckoo hashing as a Ph.D student in 2001





# Motivation of Cuckoo Hashing

- Traditional methods  $O(n)$  worst case behavior
- Suppose that all data can fit in to memory
- Can search/find/contains  $O(1)$  worst case ?
- How to accelerate the search/find/contains operations ?

# Intuition of Cuckoo Hashing

- Each item could hash to  $K$  positions /buckets
- When no position vacant, just rearrange the elements

# How Cuckoo Hashing Works

- K hash functions and K tables. (Usually  $K = 2$ )
- Insert(x)

$\text{index1} = \text{hashfunc1}(x)$   $\text{index2} = \text{hashfunc2}(x)$

if (table1[index1] vacant) put x there

else if (table2[index2] vacant) put x there

else rearrange

# Insert(99) in Cuckoo Hashing

index	key
0	12
1	34
2	56
3	vacant
4	78

Index	key
0	11
1	22
2	vacant
3	33
4	44

Hashfunc1(99) = 2

Hashfunc2(99) = 3

# Insert(99) in Cuckoo Hashing

index	key
0	12
1	34
2	56
3	vacant
4	78

Index	key
0	11
1	22
2	vacant
3	99
4	44

33

Hashfunc1(99) = 2

Hashfunc2(99) = 3

Insert 99 to table2

# Insert(99) in Cuckoo Hashing

index	key
0	12
1	34
2	56
3	vacant
4	33

Index	key
0	11
1	22
2	vacant
3	99
4	44

78

Hashfunc1(99) = 2  
Hashfunc2(99) = 3  
Insert 99 to table2  
Hashfunc1(33) = 4  
Insert 33 to table1

# Insert(99) in Cuckoo Hashing

index	key
0	12
1	34
2	56
3	vacant
4	33

Index	key
0	11
1	22
2	78
3	99
4	44

Hashfunc1(99) = 2

Hashfunc2(99) = 3

Insert 99 to table2

Hashfunc1(33) = 4

Insert 33 to table1

Hashfunc2(78) = 2

Insert 78 to table2



# Why rearrangement may fail



# Why rearrangement may fail

- Table is full
- Cycle

# What to do when insertion failed

- Change hash functions OR Resize the tables
- When  $K = 2$  and load factor  $< 50\%$

the probability of a cycle is very low

expected number of displacements is a small constant

- What about  $K = 3$  or larger ? Open Question !

# Efficiency of Cuckoo Hashing

- Insert: Amortized (Expected) const time.
- Find:  $O(1)$  worst case
- Delete:  $O(1)$  worst case
- 50% space efficient for  $K = 2$

# Implementation of Cuckoo Hashing

- Two tables is not a must. It just made the analysis much easier
- Only one table & two hash functions is enough

# Implementation of Cuckoo Hashing

- Two tables is not a must. It just made the analysis much easier
- In practice, we can use  $K$  other than 2 hash functions to be more space efficient

# Trick1 of implementation of Cuckoo hashing : d-ary cuckoo hashing

- K tables and K hash functions
- K possible positions for each key
- 91%+ space efficient
- Well known as d-ary cuckoo hashing

# Trick2 of implementation of Cuckoo hashing : bucketized cuckoo hashing

- Each location is a bucket capable of holding  $B$  entries

index	K1	K2	K3
0	12	44	vacant
1	33	30	66
2	56	40	vacant
3	11	15	37



## Trick3 of implementation of Cuckoo hashing : d-ary + bucketized cuckoo hashing

- Each key is hashed according to  $K$  hash functions, leading to  $K$  possible locations for that key
- Each location is a bucket capable of holding  $B$  entries

index	K1	K2	K3
0	12	44	vacant
1	33	30	66
2	56	40	vacant
3	11	15	37

## Trick3 of implementation of Cuckoo hashing : d-ary + bucketized cuckoo hashing

- How to choose  $K$  and  $B$  ?
- In practice ,  $K = 3$  or  $4$  will be enough
- $B * \text{sizeof}(\text{Key})$  can fit in to a cacheline

## Trick4 of implementation of Cuckoo hashing : d-ary + bucketized + vectorized cuckoo hashing

- Each entry consists of key and payload
- B Keys are stored contiguously , followed by B contiguous payloads
- Use SIMD to compare keys

index	K1	K2	K3	K4	P1	P2	P3	P4
0	11	44	77	vacant	Bob	Alice	Marry	vacant
1	22	55	vacant	vacant	Lucy	Jack	vacant	vacant
2	33	66	88	vacant	Jeff	Jure	David	vacant

# Why named Cuckoo Hashing



# Outline


- Why Separate Chaining & Linear Probe are not enough
- Cuckoo Hashing
- **Hopscotch Hashing**

# Hopscotch Hashing

- Motivation
- Intuition
- How it works
- Quiz

# Authors of Hopscotch Hashing

- **Nir Shavit**, Professor of MIT
- describe Hopscotch Hashing in 2008
- Co-author of The art of Multicore Programming







# Motivation Of Hopscotch Hashing

## Disadvantages of Cuckoo Hashing

- need to access sequences of unrelated locations on different cache lines
- need to keep the load factor smaller than 50%

# Intuition Of Hopscotch Hashing

- Each key can be located in  $H$  positions, whose elements can be fit in to a cacheline
- Element will be evicted to the new location whose distance from the original position is at most  $H$

# Hopscotch Hashing

- Motivation
- Intuition
- How it works
- Quiz

# How it works

- How to add item  $x$  where  $h(x) = i$ :
  - Starting at  $i$ , use linear probing to find an empty entry at index  $k$ .
  - If the empty entry's index  $k$  is within  $H - 1$  of  $i$ , place  $x$  there and return.
  - Otherwise, find a closer position via HopInfo

# What is HopInfo

- Each position  $i$  contains HopInfo which tells whether the item in the alternate position is occupied by an element that hashes to  $i$
- $\text{Hop}[8] = 0010$  indicates that only position 10 currently contains items whose hash value is 8, while positions 8, 9, and 11 do not.

index	item	Hop
6	G	1000
7	H	1100
8	M	0010

# How to find a closer position

- Otherwise, find a closer position via HopInfo

To create an empty entry closer to  $i$ , find an item  $y$  whose hash value  $m$  lies between  $i$  and  $k$ .

$m$  satisfies that  $k - H < m < k$

Displacing  $y$  to  $k$  creates a new empty slot closer to  $i$ . Repeat. If no such item exists, or if the bucket already  $i$  contains  $H$  items, resize and rehash the table.

# Insert(K)

index	item	HopInfo
6	C	1000
7	A	1100
8	D	0010
9	B	1010
10	E	0000
11	H	0010
12	F	1000
13	G	0000
14	vacant	0000

A:7 B:9 C:6  
D:7 E:8 F: 12  
G: 11 H: 9 K: 6

H = 4

# Insert(K)

index	item	HopInfo
6	C	1000
7	A	1100
8	D	0010
9	B	1010
10	E	0000
11	H	0010
12	F	1000
13	G	0000
14	vacant	0000

A:7 B:9 C:6  
D:7 E:8 F: 12  
G: 11 H: 9 K: 6

Linear Probing suggests 14

$$14 - 6 \geq H$$

TOO FAR !



# Insert(K)

index	item	HopInfo
6	C	1000
7	A	1100
8	D	0010
9	B	1010
10	E	0000
11	H	0010
12	F	1000
13	G	0000
14	vacant	0000

A:7 B:9 C:6  
D:7 E:8 F: 12  
G: 11 H: 9 K: 6

Linear Probing suggests 14

$$14 - 6 \geq H$$

TOO FAR !

Consult HopInfo[11]  
Move G Down

# Insert(K)

index	item	HopInfo
6	C	1000
7	A	1100
8	D	0010
9	B	1010
10	E	0000
11	H	0001
12	F	1000
13	vacant	0000
14	G	0000

A:7 B:9 C:6  
D:7 E:8 F: 12  
G: 11 H: 9 K: 6

Linear Probing suggests 14

$$14 - 6 \geq H$$

TOO FAR !

Consult HopInfo[11]  
Move G Down

# Insert(K)

index	item	HopInfo
6	C	1000
7	A	1100
8	D	0010
9	B	1010
10	E	0000
11	H	0001
12	F	1000
13	vacant	0000
14	G	0000

A:7 B:9 C:6  
D:7 E:8 F: 12  
G: 11 H: 9 K: 6

$13 - 6 \geq H$

Still TOO FAR !

Consult HopInfo[10] Does Not Help

Consult HopInfo[11] Does Not Help

Consult HopInfo[12] Move F Down

# Insert(K)

index	item	HopInfo
6	C	1000
7	A	1100
8	D	0010
9	B	1010
10	E	0000
11	H	0001
12	vacant	0100
13	F	0000
14	G	0000

A:7 B:9 C:6  
D:7 E:8 F: 12  
G: 11 H: 9 K: 6

$13 - 6 \geq H$

Still TOO FAR !

Consult HopInfo[10] Does Not Help

Consult HopInfo[11] Does Not Help

Consult HopInfo[12] Move F Down

# Insert(K)

index	item	HopInfo
6	C	1000
7	A	1100
8	D	0010
9	B	1010
10	E	0000
11	H	0001
12	vacant	0100
13	F	0000
14	G	0000

A:7 B:9 C:6  
D:7 E:8 F: 12  
G: 11 H: 9 K: 6

$12 - 6 \geq H$   
Still TOO FAR  
Consult HopInfo[9] Move B Down

# Insert(K)

index	item	HopInfo
6	C	1000
7	A	1100
8	D	0010
9	vacant	0011
10	E	0000
11	H	0001
12	B	0100
13	F	0000
14	G	0000

A:7 B:9 C:6  
D:7 E:8 F: 12  
G: 11 H: 9 K: 6

$12 - 6 \geq H$   
Still TOO FAR  
Consult HopInfo[9] Move B Down

# Insert(K)

index	item	HopInfo
6	C	1000
7	A	1100
8	D	0010
9	vacant	0011
10	E	0000
11	H	0001
12	B	0100
13	F	0000
14	G	0000

A:7 B:9 C:6  
D:7 E:8 F: 12  
G: 11 H: 9 K: 6

$9 - 6 < H$   
GOOD ! Insert K there

# Insert(K)

index	item	HopInfo
6	C	1001
7	A	1100
8	D	0010
9	K	0011
10	E	0000
11	H	0001
12	B	0100
13	F	0000
14	G	0000

A:7 B:9 C:6  
D:7 E:8 F: 12  
G: 11 H: 9 K: 6

$9 - 6 < H$   
GOOD ! Insert K there



## Find(K) and Del(K)

- $\text{index} = \text{hash}(K)$ . Just check positions range from index to  $\text{index} + H - 1$  to locate or / and del it .
- Both Find and Del cost  $O(1)$  time
- Insert costs expected const time

# Hopscotch Hashing

- Motivation
- Intuition
- How it works
- Quiz

# Quiz

- What's the difference between CH and HH
- Is the sequence of displacements possible be cyclic in HH
- When we have to resize in HH
- Will HH work without hopinfo
- What data structures can be used to store hopinfo

# Next Episode Preview

- BloomFilter
- HyperLogLog



# Thank You All